

Stochastic Gradient Descent AdaGrad ADAM

M. Fatih Amasyalı

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Finite sum of functions

- In ML optimization, objective function has the form:

$$f(x) = \frac{1}{n} \sum_{i=1}^n f_i(x)$$

n : training set size

$f_i(x)$: loss value for i .th training point according to parameter x

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Some examples of finite sum

- Training set: $i=1:n$ a_i : i.th input, b_i : i.th output
- x : parameters of a model to be trained
- Least squares:

$$\min_x \frac{1}{n} \sum_{i=1}^n (a_i^\top x - b_i)^2 = \min_x \frac{1}{n} \sum_{i=1}^n (f_i(x))^2$$

- Lasso:


$$\min_x \frac{1}{n} \sum_{i=1}^n (a_i^\top x - b_i)^2 + \lambda \sum_{j=1}^d |x_j|$$

- ANN:

$$\min_x \frac{1}{n} \sum_{i=1}^n \text{loss}(\text{ANN}(x, a_i), b_i)$$

Gradient Descent

- $x_{t+1} = x_t - \eta_t \nabla f(x_t) = x_t - \eta_t \sum_{i=1}^n \nabla f_i(x_t)$
- $f(x_1, x_2) = x_1^2 + 2x_2^2$
- Learning rate (η_t) = 0.1
- it converges at 40.th step
- Code: `hessian_gradyan.m`



Large scale problems


- Zillions of parametres (d)
- Zillions of data points (n)
- Each GD iteration requires $n \cdot d$ calculations

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes


YILDIZ TECHNICAL UNIVERSITY

A simple model

- $\min_x \frac{1}{2n} \sum_{i=1}^n (a_i^T x - b_i)^2 = \min_x \frac{1}{2n} \sum_{i=1}^n (f_i(x))^2$
- a_i, b_i, x : scalar (1 dim.), $n=8$
- Code: `sgd_how.m`



Training points, - true model $b=1a$,
-- current model $b=-10a$




Corresponding $f_i(x)$ s = $(a_i x - b_i)^2$

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

A simple model



Obverse that:

For all i , $\nabla f_i(x)$ and $\nabla f(x)$ have the same sign outside R

So, if we use $\nabla f_i(x)$, we guarantee the improvement outside R

In R , we get chaos ☺

Lets see

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Stochastic Gradient Descent

- $x_{t+1} = x_t - \eta_t \nabla f_i(x_t)$
- $f(x_1, x_2) = x_1^2 + 2x_2^2$
- Learning rate (η_t) = 0.1
- it does not converge
- Code: `sgd.m`

Region R : an area (in 2 dim)

Outside R , GD and SGD are very similar


Inside R (near x^*), SGD fluctuates

Usage of $\text{mean}(x_t)$ or $\text{mean}(x_{t-k} \dots x_t)$ instead of the last x_t make sense

Here, the last $x_t = [-0.2692, 0.3260]$, $\text{mean}(x_t) = [-0.1192, 0.0452]$

Each GD iteration requires $d \cdot n$ calculations

Each SGD iteration requires d calculations ☺



Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Stabilization of SGD

- Gradient clipping[*]: rescale gradients so that their norm is at most a predetermined value (c).

if $\|\nabla f_i(x_t)\| \geq c$ then $\nabla f_i(x_t) = c \frac{\nabla f_i(x_t)}{\|\nabla f_i(x_t)\|}$

- Mini-batch: instead of one example, use k example. Reduce variance of $\nabla f_i(x_t)$

$$x_{t+1} = x_t - \frac{\eta_t}{|K_t|} \sum_{j \in K_t} \nabla f_j(x_t)$$


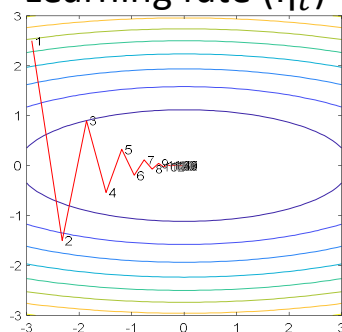
[*] <https://towardsdatascience.com/what-is-gradient-clipping-b8e815cdfb48>

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Gradient Descent

- $f(x_1, x_2) = x_1^2 + 8x_2^2$
- Learning rate (η_t) = 0.1



GD wastes time. 2.order information helps but requires $d \times d$ calculations
There is an another way, mimic Hessian (quasi-Newton methods)
Observe all previous gradients

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

AdaGrad[*]: A quasi-Newton method

- $G_t = \sum_{j=1}^t \nabla f(x_j) \nabla f(x_j)^\top$
- G_t (sum of outer product of all previous gradients) mimic Hessian
- G_t : $d \times d$ matrix
- 2 versions:
 - $x_{t+1} = x_t - \eta_t G_t^{-1/2} \nabla f(x_t)$
Full matrix inversion and square root requires $d \times d$ calculations
 - $x_{t+1} = x_t - \eta_t \text{diag}(G_t)^{-1/2} \nabla f(x_t)$
Diagonal matrix (vector in d dim) inversion and square root requires d calculations

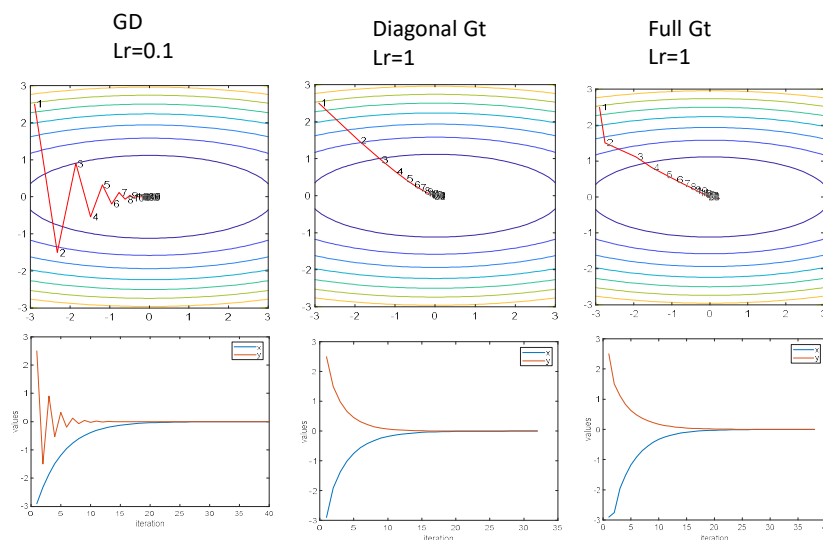
[*] <http://www.jmlr.org/papers/volume12/duchi11a/duchi11a.pdf>

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

code: hessian_gradyan.m

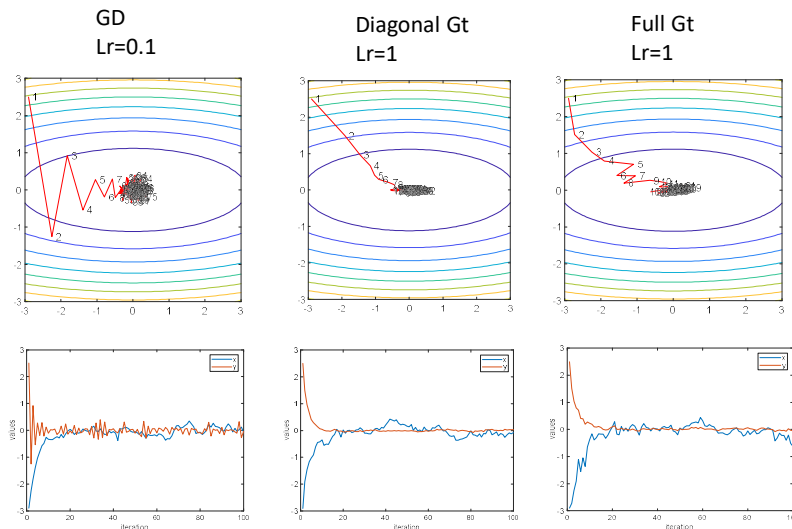
AdaGrad



Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Stochastic AdaGrad



Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

ADAM[*]

Algorithm 1: *Adam*, our proposed algorithm for stochastic optimization. See section 2 for details, and for a slightly more efficient (but less clear) order of computation. g_t^2 indicates the elementwise square $g_t \odot g_t$. Good default settings for the tested machine learning problems are $\alpha = 0.001$, $\beta_1 = 0.9$, $\beta_2 = 0.999$ and $\epsilon = 10^{-8}$. All operations on vectors are element-wise. With β_1^t and β_2^t we denote β_1 and β_2 to the power t .

Require: α : Stepsize

Require: $\beta_1, \beta_2 \in [0, 1)$: Exponential decay rates for the moment estimates

Require: $f(\theta)$: Stochastic objective function with parameters θ

Require: θ_0 : Initial parameter vector

$m_0 \leftarrow 0$ (Initialize 1st moment vector)

$v_0 \leftarrow 0$ (Initialize 2nd moment vector)

$t \leftarrow 0$ (Initialize timestep)

while θ_t not converged **do**

$t \leftarrow t + 1$

$g_t \leftarrow \nabla_{\theta} f_t(\theta_{t-1})$ (Get gradients w.r.t. stochastic objective at timestep t)

$m_t \leftarrow \beta_1 \cdot m_{t-1} + (1 - \beta_1) \cdot g_t$ (Update biased first moment estimate)

$v_t \leftarrow \beta_2 \cdot v_{t-1} + (1 - \beta_2) \cdot g_t^2$ (Update biased second raw moment estimate)

$\hat{m}_t \leftarrow m_t / (1 - \beta_1^t)$ (Compute bias-corrected first moment estimate)

$\hat{v}_t \leftarrow v_t / (1 - \beta_2^t)$ (Compute bias-corrected second raw moment estimate)

$\theta_t \leftarrow \theta_{t-1} - \alpha \cdot \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$ (Update parameters)

end while

return θ_t (Resulting parameters)

[*] Kingma, D. P., & Ba, J. Adam: A method for stochastic optimization. ICLR 2015, *arXiv:1412.6980*.

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

code: moving_avg.m

ADAM

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

$$\text{Adagrad version of } v_t: v_t = v_{t-1} + g_t^2$$

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

$$\text{Adagrad version of } m_t: m_t = g_t$$


$$\hat{m}_t = m_t / (1 - \beta_1^t)$$

Corrections

$$\hat{v}_t = v_t / (1 - \beta_2^t)$$

$$\theta_t = \theta_{t-1} - \alpha \hat{m}_t / (\sqrt{\hat{v}_t} + \epsilon)$$

Update




Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Adam update

Robust to magnitude of gradients, noisy gradients (stochastic opt.), sparse gradients





Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

code: myADAM.m


ADAM in 2d



It directly goes to the optimal point.
Does it know the optimum point ☺

No ☺



Its update values are equal for all dimensions



Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Disadvantage of ADAM



When if gradient is very close to 0 for a dimension,
it can make it bigger.

Mehmet Fatih AMASYALI Optimization Techniques Lecture Notes

YILDIZ TECHNICAL UNIVERSITY

Other quasi Newton methods

- Broyden–Fletcher–Goldfarb–Shanno (BFGS)
- Davidon–Fletcher–Powell (DFP)

References

- SGD related sections:
 - Survit Sra
<https://www.youtube.com/watch?v=k3AiUhwHQ28>
 - Gilbert Strang
<https://www.youtube.com/watch?v=AeRwohPuUHQ>
- AdaGrad:
 - Duchi, J., Hazan, E., & Singer, Y. (2011). Adaptive subgradient methods for online learning and stochastic optimization. *Journal of machine learning research*, 12(7).
<http://www.jmlr.org/papers/volume12/duchi11a/duchi11a.pdf>